

Chapter 2:

System Structures

Chien Chin Chen

Department of Information Management
National Taiwan University

Objectives

- To describe the **services** an operating system provides to users, processes of the system.
- To discuss the various ways of **structuring an operating system**.
- To explain how operating systems are installed and customized and how they boot.

Operating System Services (1/4)

- One set of operating-system services provides functions that are **helpful to the user** (or user processes):
 - **User interface** - almost all operating systems have a user interface (UI).
 - **Command-line interface** (CLI): require a program to allow entering and editing of text commands.
 - **Graphics user interface** (GUI): a window system with a pointing device and a keyboard to enter commands.
 - **Batch**: commands and directives are entered into files to be executed.
 - **Program execution** - the system must be able to load a program into memory and to run that program, end execution, either normally or abnormally.

Operating System Services (2/4)

- **I/O operations** - a user program may require I/O.
 - For efficiency and protection, users cannot control I/O devices directly.
 - The operating system must provide a means to do I/O.
- **File-system manipulation** - user programs need to read/write/create/delete/search files and directories.
 - The operating system provides permission management to allow or deny access to files or directories.
- **Communications** – user processes may exchange information, on the same computer or between computers over a network.
 - Communications may be via *shared memory* or through *message passing*.
- **Error detection** – the operating system needs to be constantly aware of possible errors.
 - And fix errors generated from hardware (disk fail) or software (arithmetic error).

Operating System Services (3/4)

- For systems with **multiple users** (processes), another set of operating-system functions exists for ensuring the efficient operation of the system itself.
 - **Resource allocation** - when multiple users or multiple jobs running concurrently, resources must be allocated to each of them.
 - CPU, memory, file storage ...
 - Operating systems have *CPU-scheduling routines* to determine the best way to use the CPU.
 - **Accounting** - To keep track of which users use how much and what kinds of computer resources.
 - Usage statistics may be a valuable tool for researchers who wish to reconfigure the system to improve computing services.

Operating System Services (4/4)

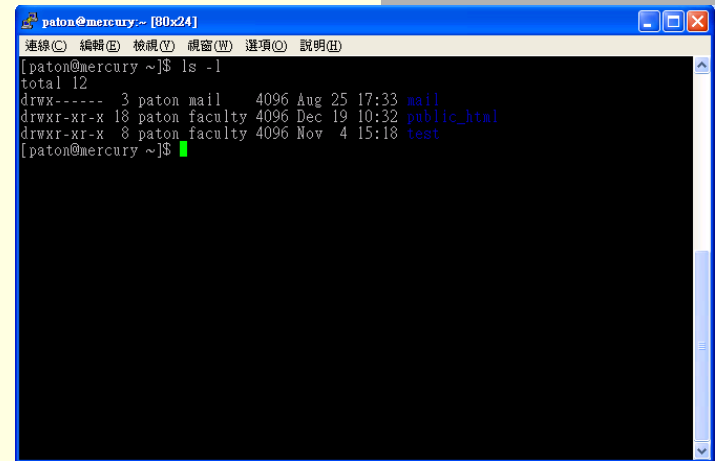
- **Protection and security** - a multiuser or networked computer system may want to control use of user information.
 - **Concurrent** processes should not interfere with each other.
 - **Protection** involves ensuring that all access to system resources is controlled.
 - **Security** of the system from outsiders requires user authentication, extends to defending external I/O devices (e.g., network adapters) from invalid access attempts.

User Operating-System Interface – **CLI** (1/4)

■ *Command-line interface*

(interpreter):

- Primarily **get** and **execute** the next **user-specified command**.
- Many of the commands are to manipulate files/directories.
 - Create/delete/list/print/copy/execute...



```
paton@mercury:~ [80x24]
連線(C) 編輯(E) 檢視(V) 选项(O) 说明(H)
[paton@mercury ~]$ ls -l
total 12
drwx----- 3 paton mail 4096 Aug 25 17:33 mail
drwxr-xr-x 18 paton faculty 4096 Dec 19 10:32 public_html
drwxr-xr-x 8 paton faculty 4096 Nov 4 15:18 test
[paton@mercury ~]$
```



```
命令提示字元
C:\>dir /w
磁碟区 C 中的磁碟是 System
磁碟区序號: A414-78C5

C:\> 的目錄

AUTOEXEC.BAT          CONFIG.SYS             [Documents and Settings]
[drivers]              drivez.log             [1386]
[Intell]               lnab.log               [Program Files]
[SUPPORT]              [SWHARE]               [SVTOOLS]
syslevel.lgl          [temp]                 TPHKLOCK.TXT
[VALUEADD]             [WINDOWS]

        6 個檔案          2,855 位元組
       11 個目錄       5,928,386,560 位元組可用

C:\>
```

User Operating-System Interface – CLI (2/4)

- Two ways to implement commands and command interpreters:
 - The command interpreter contains the code to execute the command.
 - For example, a command to delete a file may cause the interpreter to jump to a section of its code that makes the appropriate system call.
 - The number of commands that can be given determines the size of the interpreter.
 - An alternative approach implements most commands through system programs.
 - If the interpreter does not understand the command ...
 - It **identify** the command file and **load** it into memory for **execution**. 在特定的目錄下去讀指令

User Operating-System Interface – CLI (3/4)

- (cont.)
 - An UNIX example: `rm file.txt`
 - The interpreter search for a file called `rm` (`/bin/rm`).
 - **Load** it into memory and **execute** it with the parameter `file.txt`.
 - The function `rm` is completely defined by the code in the file `/bin/rm`.
 - Programmers can add new commands to the system easily.
 - The interpreter program can be **small**, and does not have to be changed for new commands to be added.
 - Used mostly among operating system, e.g., UNIX.

User Operating-System Interface – CLI (4/4)

- CLI is sometimes implemented in kernel, sometimes by systems program.
- An operating system can have multiple 獨立的program interpreters to choose from, known as **shells**.
 - For example, on UNIX and Linux systems, there are *Bourne/C/Korn...shell*.
 - The name 'shell' originates from shells being an outer layer of interface between the user and the innards of the operating system (the kernel).
 - Most shells provide similar functionality with only minor differences; most users choose a shell based upon personal preference.
 - E.g., the syntax of **shell script**.

給shell跑的program

User Operating System Interface – **GUI** (1/2)

- A GUI provides a **desktop** metaphor interface.
 - **Icons** represent files, programs, actions, etc.
 - A **mouse click** can invoke a program, select a file ...
- GUI Timeline:
 - Experimentally appeared in the early **1970s**.
 - Became widespread by Apple Macintosh computer (Mac OS) in the 1980s.
 - Dominated by Microsoft Windows (3.1, NT, 95, 98, 2000, XP, Vista).
- UNIX systems have been dominated by command-line interface.
 - Although there are various GUI interface available.
 - X-Windows systems, K Desktop Environment (KDE) by GNU project (open source – source code is in the public domain).

User Operating System Interface – GUI (2/2)

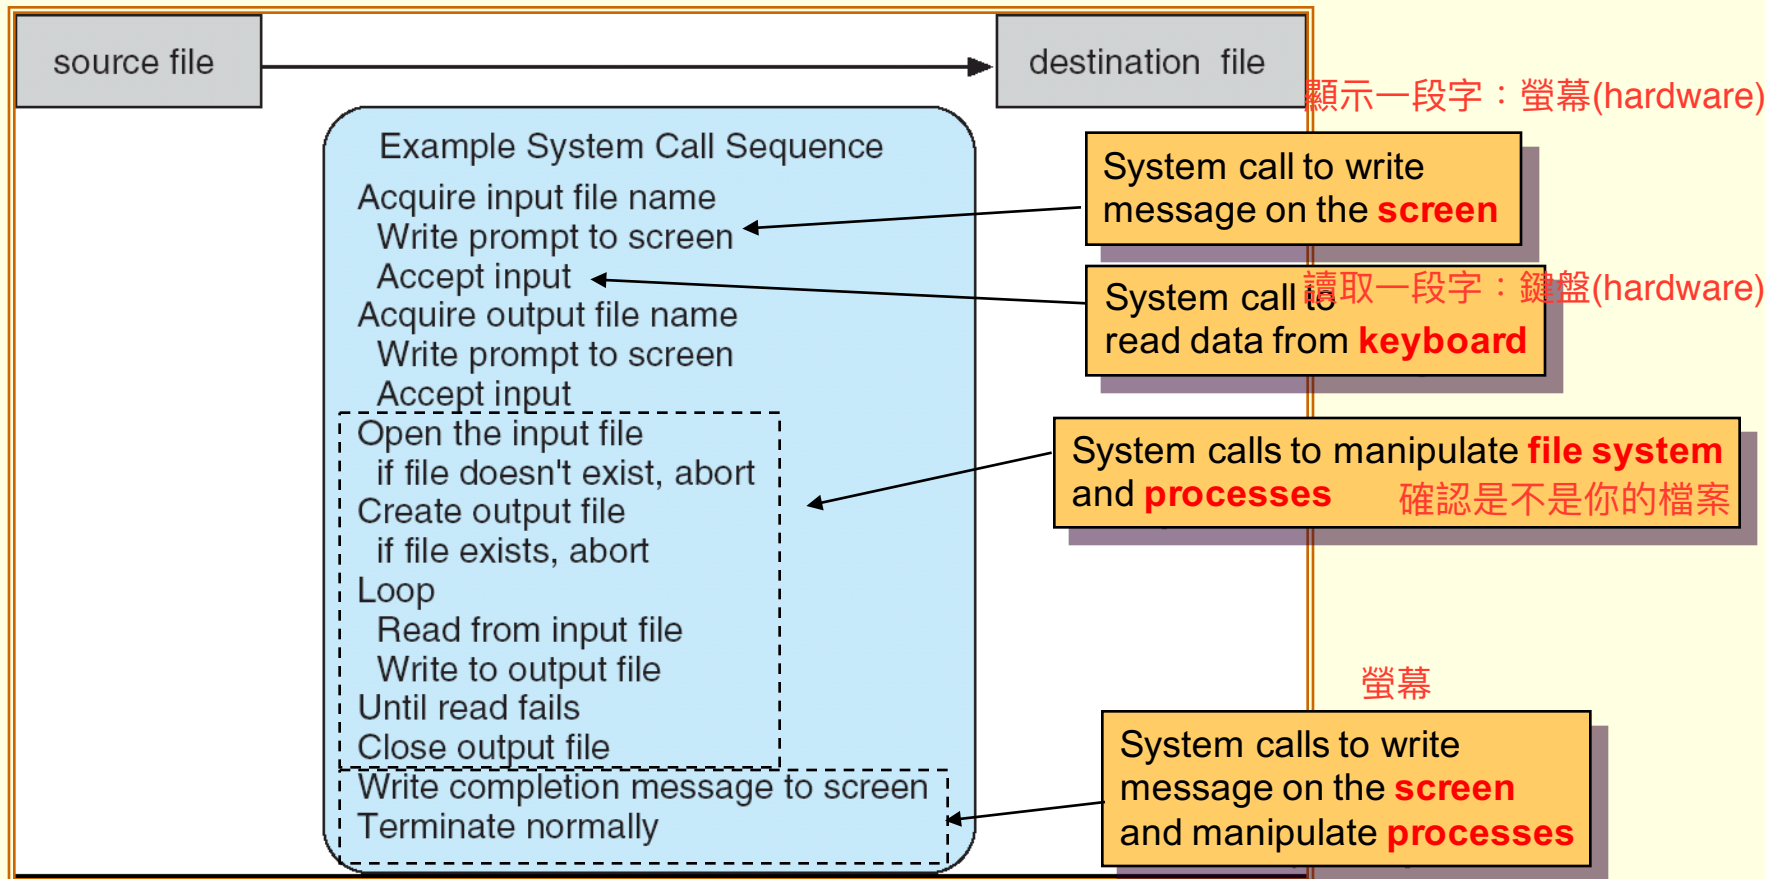
- Preference:
 - Many UNIX users prefer a command-line interface.
 - Most Windows user are pleased to use the Windows GUI and never use the MS-DOS shell interface.
- Nevertheless, many systems now include both CLI and GUI interfaces.

System Calls (1/11)

- Can be regarded as a **programming interface** to **the services provided by the OS**.
 - Called by user applications. service就是一個一個的routine,存在OS裡(kernel mode)
希望使用者寫程式叫他 (user mode)
- Are generally available as **routines**, typically written in a high-level language (C or C++).
 - Also in **low-level assembly language** (for accessing hardware).

System Calls — An Example Program of **File Copy** (2/11)

- System call sequence to copy the contents of one file to another file.



Even a simple program makes heavy use of system calls!!

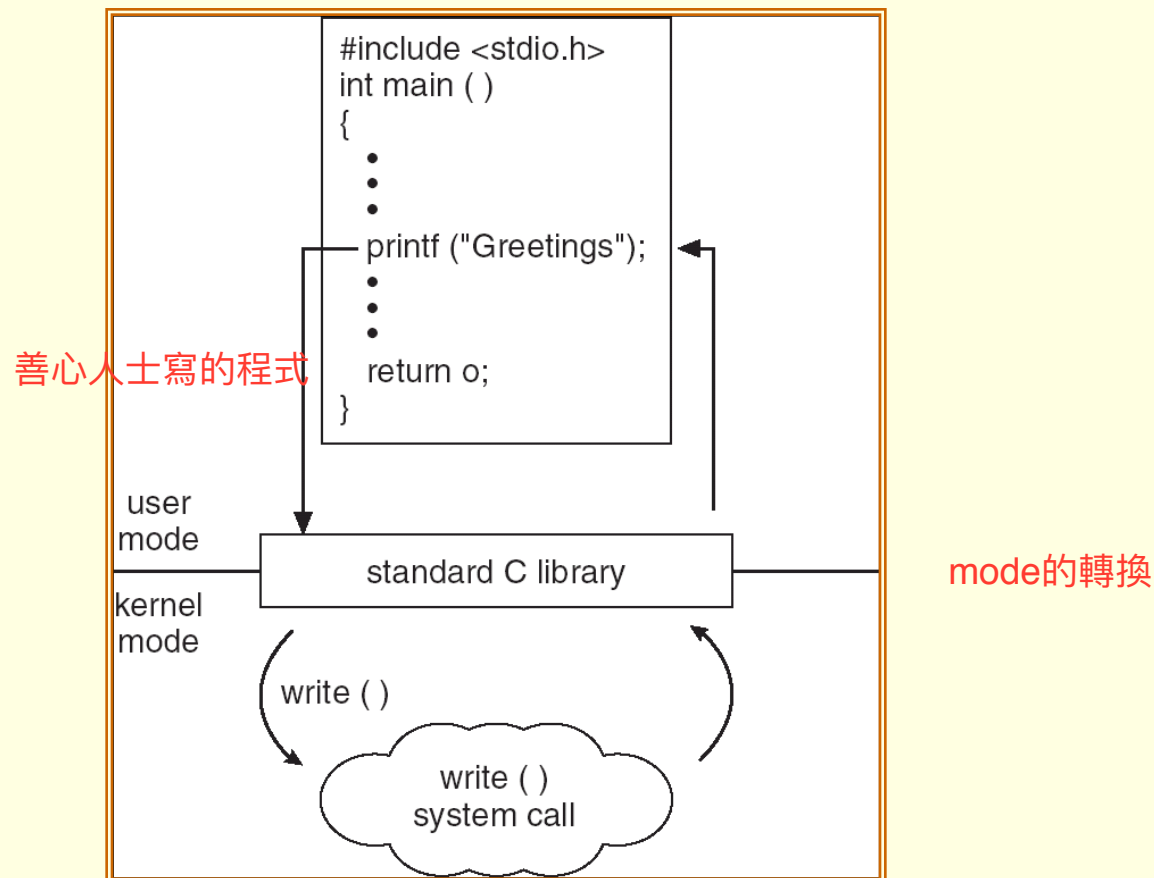
System Calls (3/11)

- Programs mostly access system services via a high-level **application program interface** (API) rather than using system call directly.
 - API specifies a set of functions (specifications) that are available to programmers.
 - **The functions of the API invoke the actual system calls on behalf of the programmer.**
- Three most common APIs:
 - **Win32 API** for Windows.
 - **POSIX API** for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X).
 - **Java API** for the Java virtual machine (JVM).

Portable Operating System Interface

System Calls (4/11)

- C program invoking `printf()` library call, which calls `write()` system call.



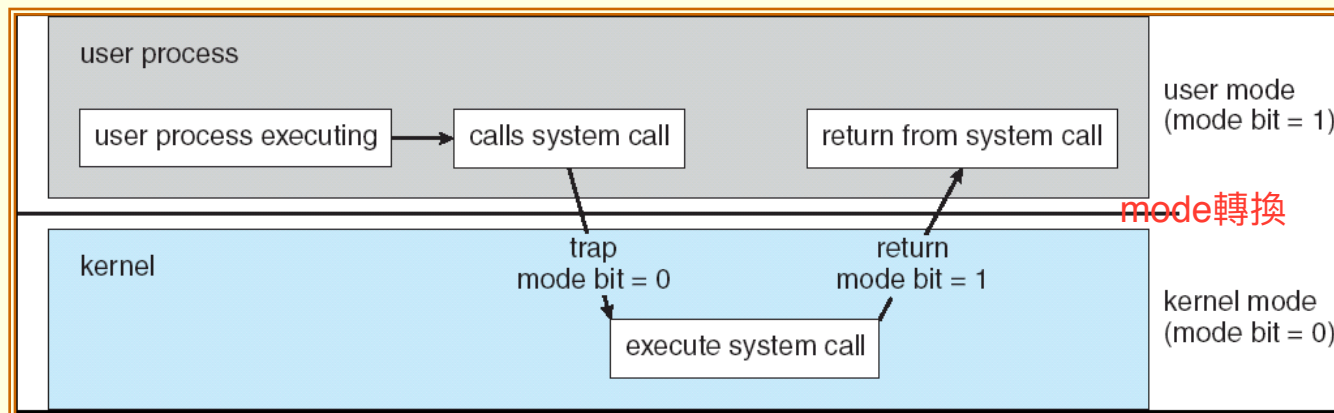
System Calls (5/11)

- Why use APIs rather than system calls?
 - **Program portability** — a program using an API can be expected to compile and run on any system that supports the same API.
 - E.g., your Windows programs on various versions of Windows.
 - Programming with API is more **simpler** than using system calls.

System Calls (6/11)

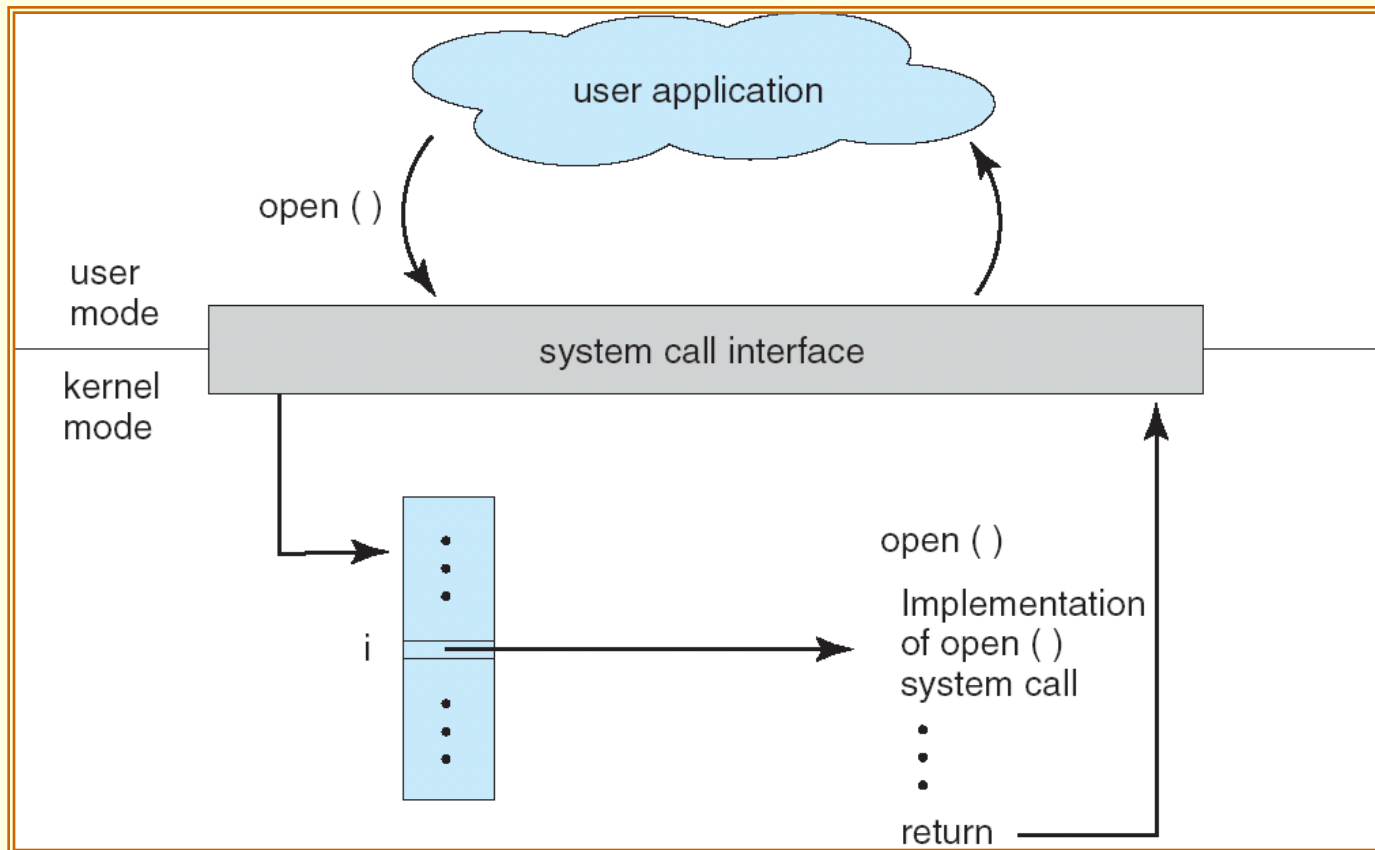
- As system calls are routines in kernel space, using it causes a **change in privileges**.
- How ?
 - Via software interrupt (e.g., `INT 0x80` assembly instruction on Intel 386 arch of Linux system).
- But before that ...
 - Similar to hardware interrupt, we need a number (index) to indicate the required system call, which is store in the `EAX` register.
- System contains a table of code pointers.
 - Using the system call number, we jump to the address of the system call for execution.

API helps us wrap all the details by simply invoking a library function.



System Calls (7/11)

- How the operating system handles a user application invoking the `open()` system call through API.



System Calls (8/11)

- To **link system call** made available by the operating system:
 - The **run-time library** for most programming languages provides a **system-call interface**.
 - Typically, a **number** associated with each system call.
 - System-call interface (or kernel) maintains a **table** indexed according to these numbers.
 - The system call interface invokes intended system call in operating system kernel and returns status of the system call and any return values.

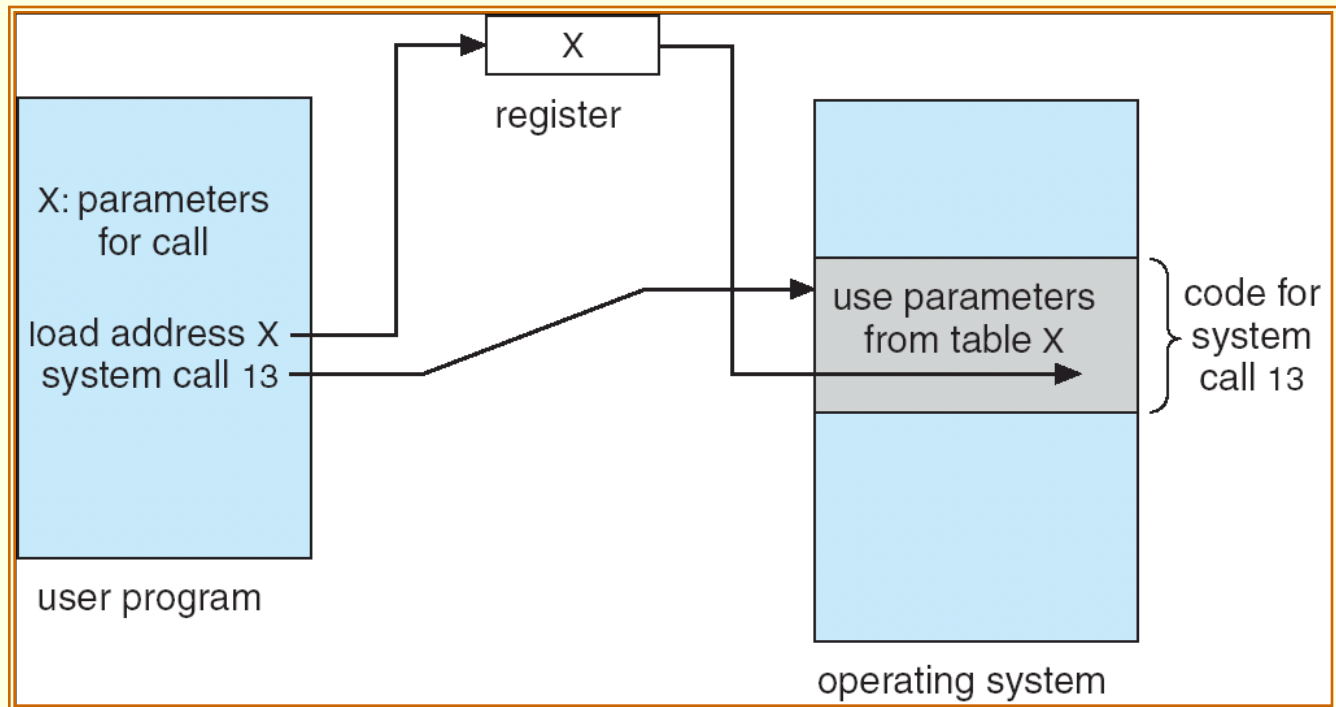
System Calls (9/11)

- With the help of API:
 - The caller need know nothing about how the system call is implemented.
 - Just needs to obey API and understand what the operating system will do as a result of that system call.
 - Details of the operating system are hidden from programmer.

System Calls (10/11)

- Often, more information is required than simply identity of desired system call
 - E.g., system call parameters.
- Three general methods used to pass parameters to the OS:
 - Simplest: pass the parameters in *registers*
 - In some cases, may be more parameters than registers.
 - Parameters are stored in a *block in memory*, and address of block passed as a parameter in a register.
 - This approach taken by Linux and Solaris.
 - Parameters are *pushed onto a stack* by the program and *popped* off the stack by the operating system.
 - Block and stack methods are popular because they do not limit the number or length of parameters being passed.

System Calls (11/11)



Types of System Calls (1/9)

- Many of today's operating system have hundreds of system calls.
 - Linux has 319 (or more) different system calls.
- Those system calls can be grouped roughly into five major categories:
 - Process control.
 - File management.
 - Device management.
 - Information maintenance.
 - Communications.

Types of System Calls —

Process Control (2/9)

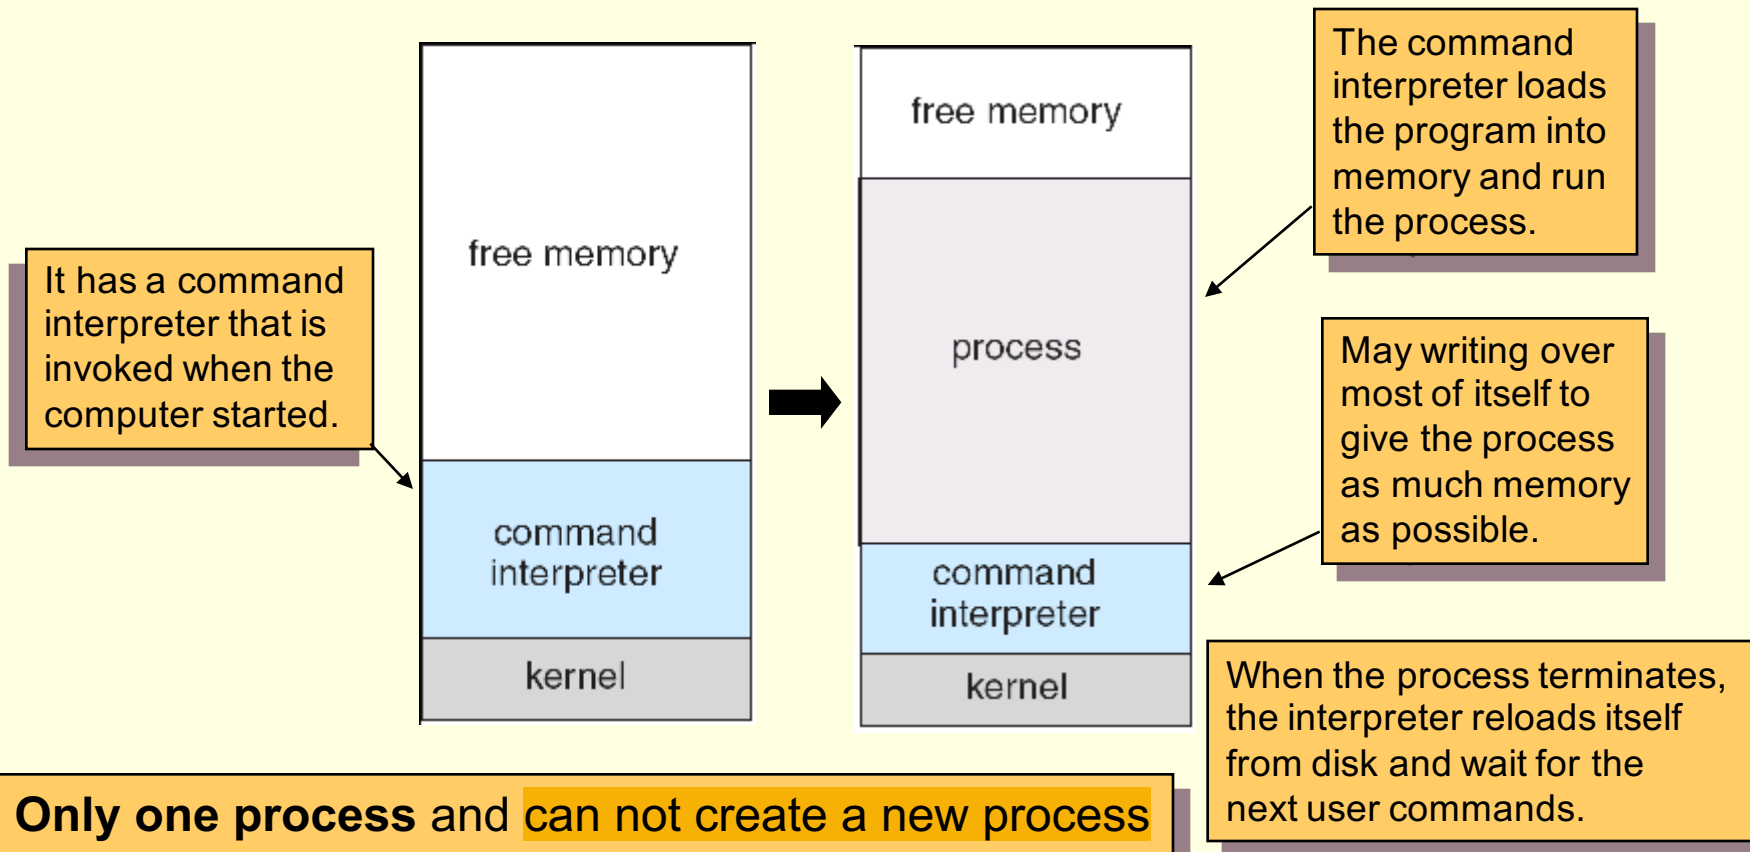
- `end and abort:`
 - A running program needs to be able to halt its execution either normally or abnormally.
 - The operating system then transfers control to the invoking command interpreter to read the next command.

- `load and execute:`
 - A process executing one program may want to load and execute another program.
 - The existing process can be lost, saved, or allowed to continue execution concurrently with the new process.
 - Chapter 6 (synchronization) discusses coordination of concurrent processes in great detail.

- `wait time/event:`
 - Having created new processes, we may need to wait for them to finish their execution.
 - We may want to wait for a certain amount of time to pass.
 - We may want to wait for a specific event to occur.

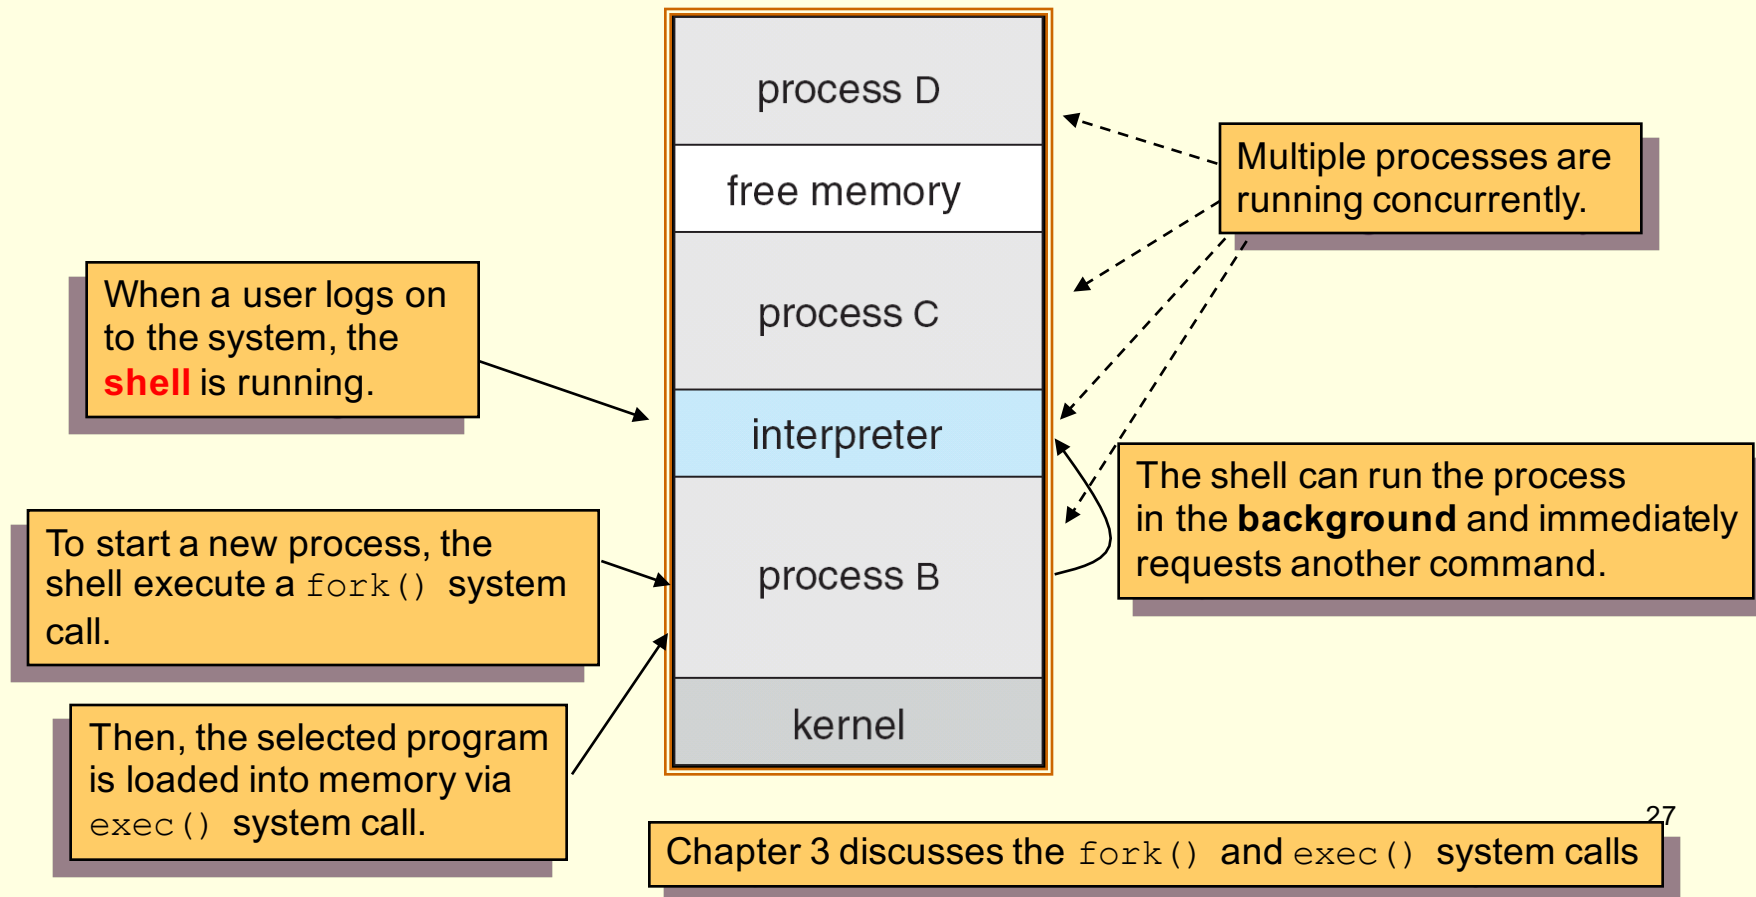
Types of System Calls — Process Control (3/9)

- Two popular variations in process control:
 - **Single-tasking system:** the MS-DOS operating system.



Types of System Calls — Process Control (4/9)

- **A multitasking system**: the FreeBSD (derived from Berkeley UNIX).



Types of System Calls — **File Management (5/9)**

- create and delete:
 - Able to create and delete files/directories.
- open and close:
 - Able to open and close a file.
- read, write, and reposition:
 - Able to read, write, or skipping to the end/head of the file.
- Other system calls for obtaining/setting file/directory attributes.

Types of System Calls — Device Management (6/9)

- The various **resources** (memory, disks, file, ...) controlled by the operating system can be thought of as **devices**.
- To access a resource, a process has to:
 - First `request` the device, to ensure **exclusive** use of it.
 - Then we can `read`, `write`, and `reposition` the device.
 - After we are finished with the device, we `release` it.
- The similarity between I/O devices and files is so great that many operating systems (UNIX) merge the two into a combined file-device structure.
 - A set of system calls is used on files and devices.
 - Sometimes, I/O devices are identified by special file names.
ex. 雷射印表機的文件
有些system call可以用在一個file or 一個device上

Types of System Calls —

Information Maintenance (7/9)

- Many system calls exist simply for the purpose of transferring information between the user program and the operating system.
 - `time` and `date` return the current time and date of the system.
 - Other system calls can return the number of current users, the amount of free memory or disk space, ...
 - Get and set processes attributes.

Types of System Calls — **Communication** (8/9)

- There are two common models of interprocess communication:
 - **Message-passing model:**
 - The communicating processes (may be on different computers) exchange messages with one another to transfer information.
 - **Client and server (daemon) architecture.**
 - Client: ask for connecting communication.
 - Server: wait for connection.
 - Require system calls to build up/terminate connection, read, and write messages.
 - **Shared-memory model:**
 - Processes use shared memory create/attach system calls to create and gain access to regions of memory owned by other processes.
 - They can then exchange information by reading and writing data in the shared areas.

Types of System Calls — Communication (9/9)

- Message passing:
 - Is useful for exchanging smaller amounts of data, because no conflicts need be avoided.
 - Is easy to implement.
- Shared memory:
 - Allows maximum speed of communication, since it can be done at memory speeds when it takes place within a computer.
 - However, problems exist in the areas of protection and synchronization between the processes sharing memory.

System Programs (1/3)

- A perspective of operating systems is a collection of **system programs**.
 - System programs provide a convenient environment for program execution and development.
 - Most users' view of the operation system is defined by system programs, not the actual system calls.
 - **Some of them are just user interfaces to system calls!!**
- Categories of system programs:
 - **File manipulation:**
 - Create, delete, copy, rename, print, dump, list, and generally manipulate files and directories

System Programs (2/3)

- **File modification:**

- Text editors to create and modify files.

- **Status information:**

- Some ask the system for info - date, time, amount of available memory, disk space, number of users
- Typically, these programs format and print the output to the terminal or other output devices

- **Programming-language support:**

- Compilers, assemblers, debuggers and interpreters sometimes provided.

- **Program loading and execution:**

- Loaders to load assembled or compiled programs into memory for execution.

System Programs (3/3)

- **Communications:**

- Provide the mechanism for creating virtual connections among processes, users, and computer systems.

- In addition to system programs, application programs are supplied to solve common problems or perform common operations.

- Web browsers, word processors, database systems, games ...

- The view of the operating system seen by most users is defined by the application and system programs, rather than the actual system calls.

Operating System Design (1/2)

- The first is to define **requirements** and **specifications**.
- However, there is no unique solution to the problem of defining the requirements for an operating system.
 - Requirements can be affected by choice of hardware, type of system.
 - Handheld devices vs. PCs.
 - Single process vs. multitasking.
- The requirement can be divided into user goals and system goals.
 - User goals – operating system should be **convenient** to use, **easy** to learn, **reliable**, **safe**, and **fast**.
 - System goals – operating system should be easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient.

Operating System Design (2/2)

- One important principle of system design is **the separation of policy from mechanism**.
 - **Policy:** **What** will be done?
 - **Mechanism:** **How** to do it?
 - Example: *timer* is a **mechanism** for ensuring CPU protection, but deciding how long the timer is to be set is a **policy** decision.
- Flexibility of the separation:
 - Policies are likely to change across places or over time.
 - The separation enables a change in policy to redefine certain policy parameters rather than changing the underlying mechanism.
- Most of Windows services mix mechanisms with policies to enforce a global look and feel.

Operating System Implementation (1/2)

- Traditionally, operating systems have been written in assembly language.
- Now, they are most written in higher-level languages such as C or C++.
 - The code can be written faster and is easier to understand and debug.
 - System is easier to port (to move to some other hardware).
 - The Linux operating system is written mostly in C and is available on a number of different CPUs, including Intel 80x86, Motorola 680X0, ...
- Previous comments on higher-level languages: reduced speed and increased storage requirements.
 - Modern compiler techniques can perform complex analysis and optimizations that produce excellent code.

Operating System Implementation (2/2)

- Moreover, major performance improvements in operating systems (and other systems) are more likely to be the result of better **data structures** and **algorithms** than of excellent assembly-language code.
- Should pay more attentions on the **memory manager** and the **CPU scheduler**.
 - They are probably the most critical routines.

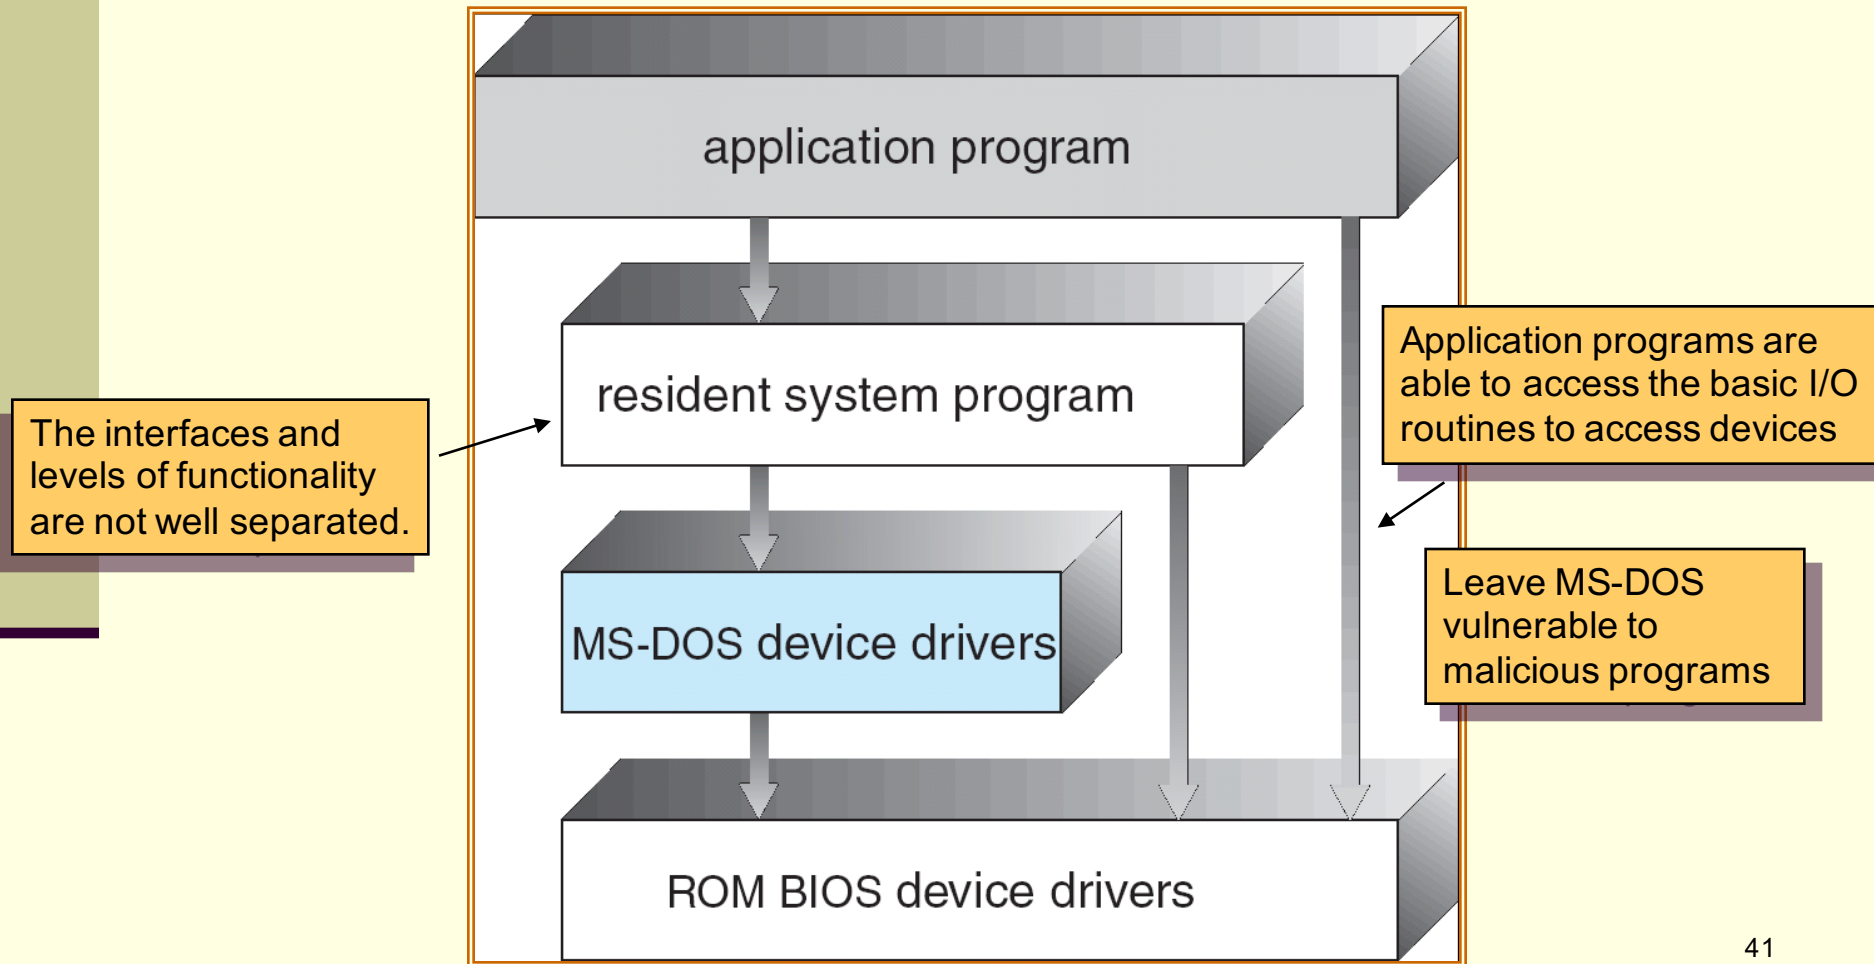
Operating-System Structure

Simple Structure (1/11)

- A common approach to implement an operating system is to partition the task into small components.
 - **Rather than have one monolithic system!!**
 - These components are interconnected and meld into a kernel.
 - But...
- **Simple structure:**
 - Many commercial system do not have well-defined structures initially.
 - Started as small, simple, and limited systems and then grew beyond their original scope.
 - For example, MS-DOS.

Operating-System Structure

Simple Structure (2/11)

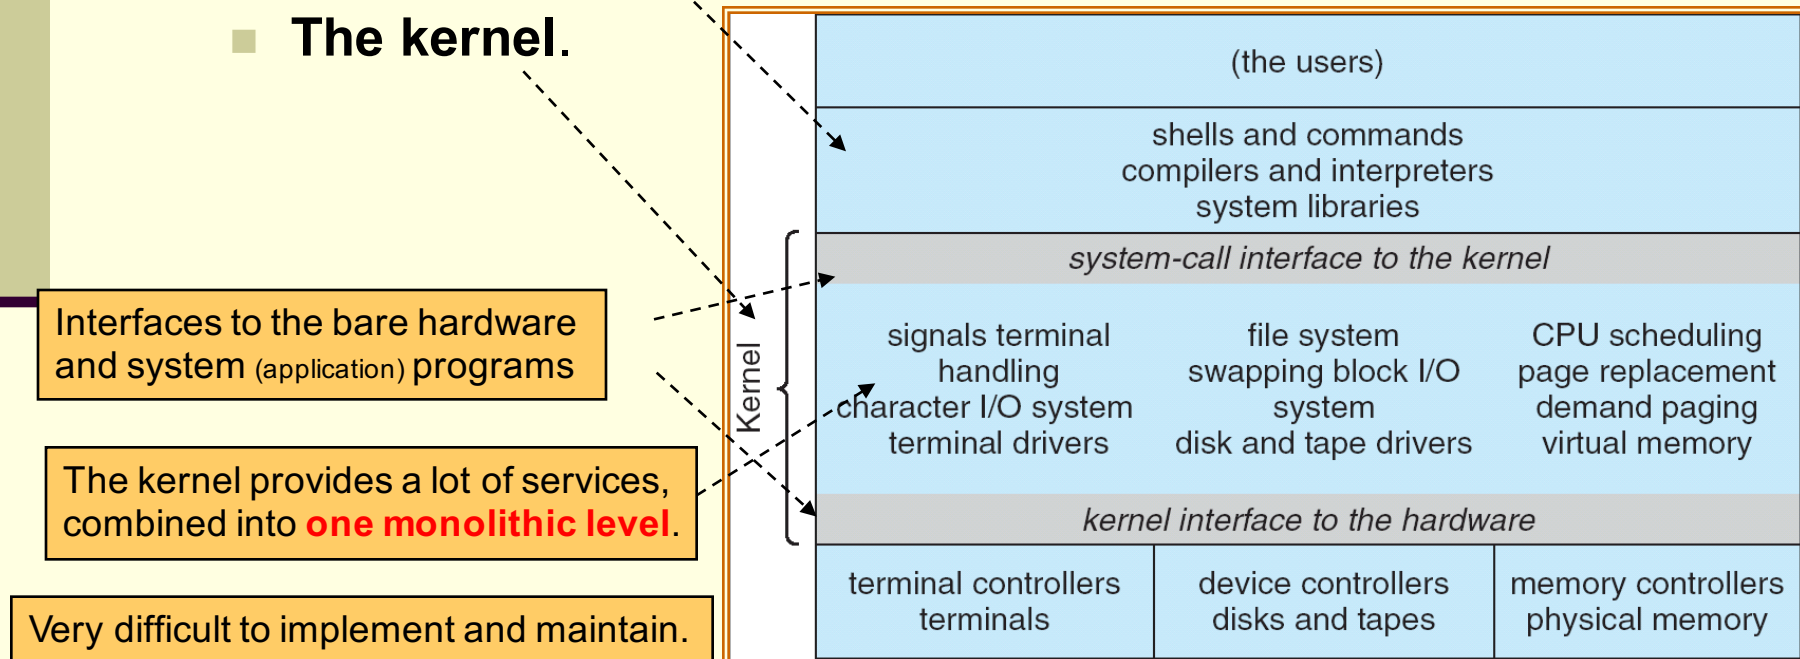


MS-DOS layer structure.

Operating-System Structure

Simple Structure (3/11)

- UNIX – limited by hardware functionality, the original UNIX operating system had limited structuring.
- The UNIX OS consists of two separable parts:
 - System programs.
 - The kernel.



Operating-System Structure

Layered Approach (4/11)


- With the improvements of hardware and programming techniques, operating systems can be broken into pieces of components.
 - That is **modular** operating systems.
 - Information hiding: hide the internal implementation detail of modules and provide external access **interfaces**.
- One way of modular system: **layered approach**.
 - The operating system is divided into a number of layers (levels), each built on top of lower layers.
 - The bottom layer (layer 0), is the hardware.
 - The highest (layer N) is the user interface.

Operating-System Structure

Layered Approach (5/11)

Layer M consists of data structures and a set of routines that can be invoked by higher-level layers.

Layer M , in turn, can invoke operations on lower-level layers.



Operating-System Structure

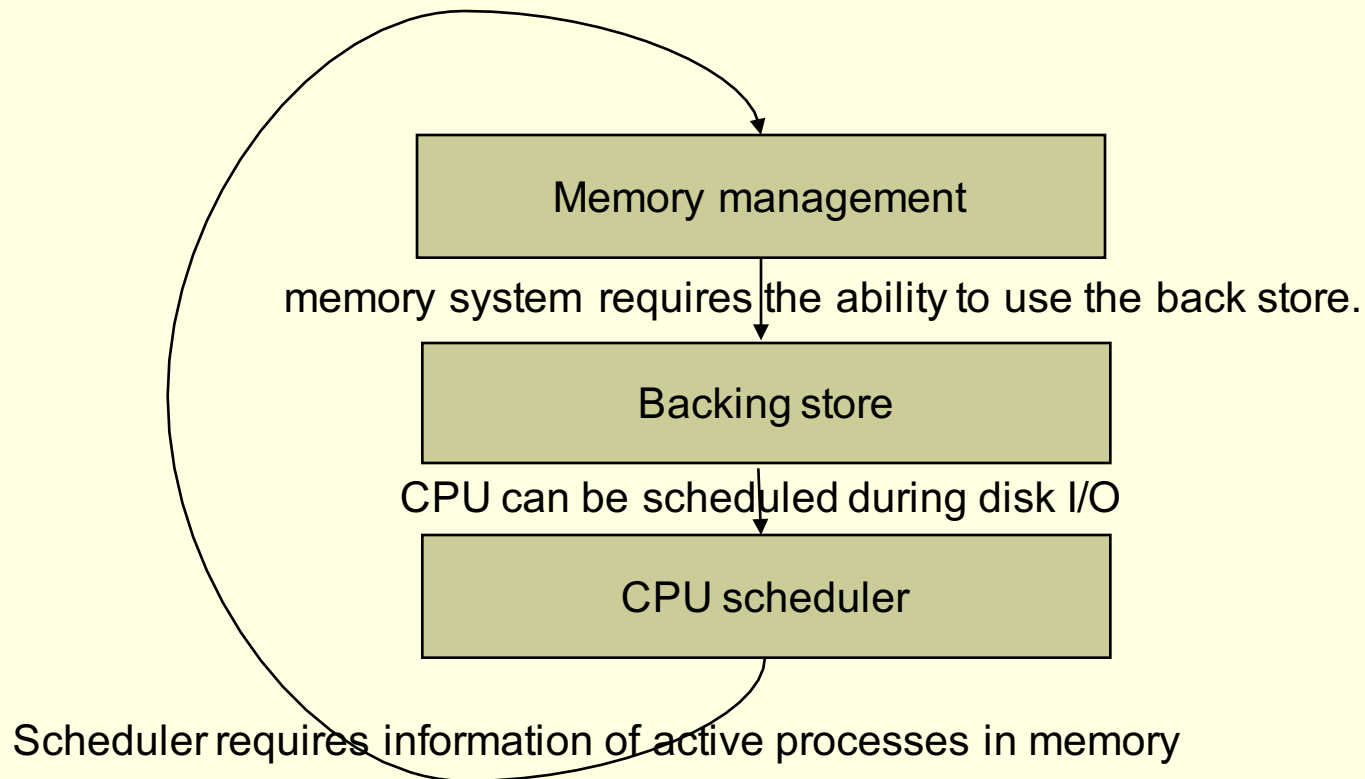
Layered Approach (6/11)

- The main **advantage** of the layered approach:
 - Simplicity of construction and **debugging**.
 - Layer-by-layer debugging, starting from layer 0.
 - If an error is found during the debugging of a particular layer, the error must be on that layer.
- The major **difficulty** of the layered approach:
 - Because only lower layers operations can be invoked, appropriately defining the various layers is difficult.
 - **System services usually tangle together.**
 - Layered implementation tend to be less efficient.
 - A function call on the top layer can lead to many lower-layer calls.
 - Function calls need to pass (redundant) parameters.
- Recently, fewer layers with more functionality are being designed.
 - Providing the advantages of modularization. 而且慢
 - Avoiding the **difficulties of layer definition and interaction.**

Operating-System Structure

Layered Approach (7/11)

■ Example of tangled layers:



Operating-System Structure

Microkernels (8/11)

- As operating systems expanded, the kernel became large and difficult to manage.
- In the mid-1980s, CMU developed an operating system called **Mach** that modularized the kernel using the **microkernel** approach.
 - Micro → removing all nonessential components from the kernel and implementing them as system and user-level programs (servers).
 - Typically, microkernels provide **process** and **memory** management, and a **communication** facility.
 - The **client program and services** communicate **indirectly** by exchanging message with the microkernel.

Operating-System Structure

Microkernels (9/11)

■ **Benefits:**

- Easier to include new operating system services to a microkernel.
 - Do not require modification of the kernel.
- The small kernel makes it easier to port to new hardware architectures.
- More **reliable and secure** (less code is running in kernel mode).

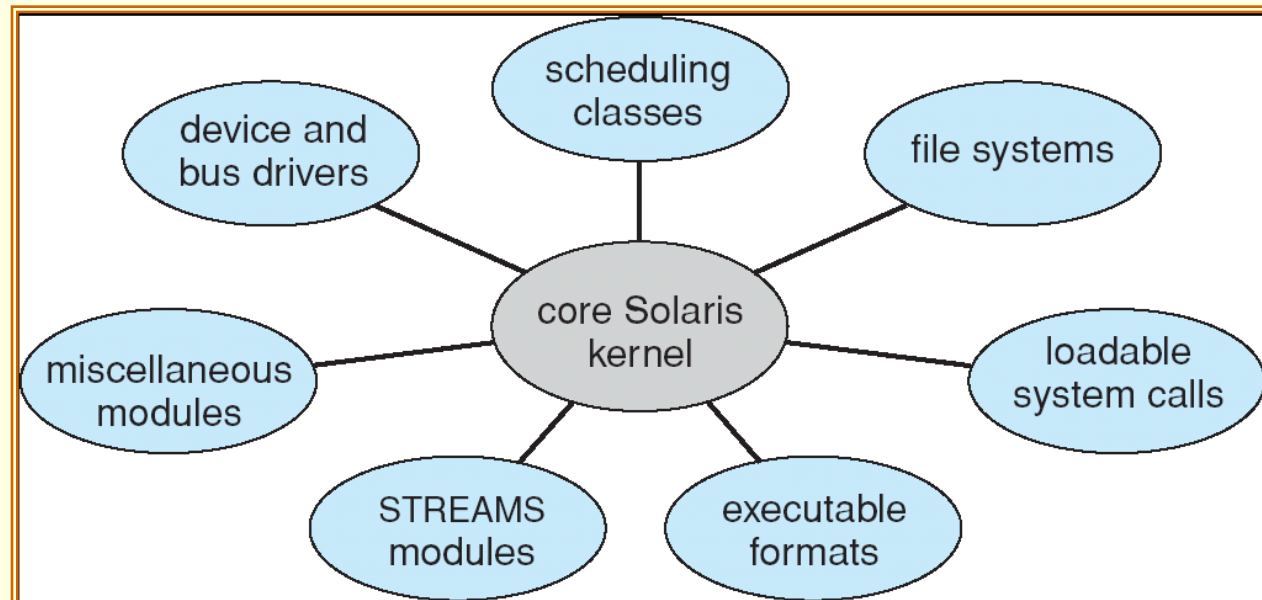
■ **Problems:** 慢到不行 (kernel->變到很多地方)

- Performance overhead of user space to kernel space communication.
- Initial Windows NT (a micorkernel organization) → Windows NT 4.0 (moving layers from user space to kernel space).

Operating-System Structure

Modules (10/11)

- A better methodology for operating-system design involves using object-oriented programming techniques to create a **modular kernel**.
 - Consists of a **core kernel**, and system service as **kernel modules**.
 - Each module talks to the others over known **interfaces**



Operating-System Structure

Modules (11/11)

- Moreover, modules (system services) can be linked into the system either during boot time or **during run time** (that is, loaded as needed within the kernel).
dynamic linking
 - Load different file system (ext2fs, FAT32 or NTFS) as needed, to save main memory.
- The module structure is similar to layered (communicate with interfaces) and microkernel approaches (a core), but with more flexible.
 - Any module can **call any other module**, but the layered approach can not.
 - Is efficient than microkernel approach because modules are in the kernel space and do not need to invoke message passing to communicate.
- The strategy of dynamically loadable modules is very popular in modern UNIX-based operating systems, such as Linux.

Operating System Generation (1/3)

- Operating systems are normally distributed on disk or CD-ROM.
 - They are designed to run on **any** class of machines **with different hardware configurations**.
- To generate a system for each specific computer site, a special program – **SYSGEM** – is needed.
 - It determines computer components by:
 - Reading a given file.
 - Asking the operator of the system for hardware information.
 - Probes the hardware directly.

Operating System Generation (2/3)

- The information must be determined:

- **CPU:**

- What CPU is to be used?
- Number of CPUs.
- Has extended instruction sets or floating point arithmetic.

- **Memory:**

- Size.

- **Devices:**

- Type and model.
- Interrupt number.

- **Operating-system options:**

- Maximum number of processes to be supported.
- CPU-scheduling algorithm.

Operating System Generation (3/3)

- Once the information is determined ...
 - Source code of the operating system can be **modified** and completely **compiled** to produce a tailored operating system.
 - System generation is slower.
 - But more specific to the underlying hardware.
 - Or, the description can cause the selection of **modules** from a **precompiled library**, which are linked together to form the operating system.
 - Because the system is not recompiled, system generation is faster.
 - The resulting system may be general.
 - Easy to modify the generated system as the hardware configuration changes (such as, add a new hardware).

System Boot (1/2)

- The generated operating system must be made available by the hardware.
- How does the hardware know where the kernel is and how to load that kernel??
- **Booting** – the procedure of starting a computer by loading the kernel.
 - Power up or reset.
- Need a **bootstrap program** to:
 - Locate the kernel on the disk.
 - Load it into memory.
 - Start its execution.
 - A simple code stored in ROM or EPROM.
 - At a fixed location so that can be loaded and executed when computer is on.
 - But before that, it first initializes all aspects of the system:
 - CPU registers, device controller, the contents of main memory ...

System Boot (2/2)

- Some computer systems (such as PCs) use a two-step booting process:
 - A simple bootstrap loader fetches a more complex boot program from disk.
 - Which in turn loads the kernel.
- The boot program stored in the **boot block** (a fixed location on disk) is usually sophisticated and modifiable and is to load an (or different) operating system into memory and begin its execution.
 - Then the operating system is said to be **running**.



End of Chapter 2

